

The Synthesis Brief

The fourteen-paper grammar of Governance as Engineering — what it claims, how strongly, and what it does not

Björn Kenneth Holmström · June 2026

Claims are labelled [R] rigorous, [IP] in-principle, or [H] heuristic.

Creative Commons Attribution-ShareAlike 4.0 International

<https://bjorkennethholmstrom.org/syntheses/brief>

A compact map of a twenty-eight-paper theoretical series, written for readers in technology and policy. It states what the framework claims, marks how strongly it claims it, reports the first prediction to be tested against data, and points to the work that has not yet been done.

What This Is, and What It Is Not

This is a synthesis of *Governance as Engineering*, a series of twenty-eight working papers that treats governance systems as feedback systems and asks what control theory, cybernetics, and information theory have to say about why they fail. The premise fits in a sentence: governance systems observe, decide, act, and observe again, and feedback systems have structural properties — latency, signal fidelity, dimensionality, gain ceilings — that bound their behaviour regardless of the intentions of the people operating them.

It is **not a manifesto**, and it is not a finished theory. The work is recent, developed by a single author working quickly and with substantial AI collaboration. Its central diagnostic claim — that many governance failures are *architectural* rather than *behavioural* — is meant to be precise enough to be wrong in identifiable ways, and therefore improvable. The point of this brief is to make the framework legible to people who might test it, break it, extend it, or build with it. A reader who wants to know whether the claims are trustworthy will look first at how they are hedged, so the hedging is placed in front rather than buried.

How to Read the Claims

Every load-bearing claim in the series carries one of three confidence labels. They are used in this brief and throughout the papers.

- **[R] Rigorous** — proven within a formal model, or a direct consequence of an established result in control theory or information theory.
- **[IP] In-principle** — the formal core is solid, but its translation into institutions is an interpretive correspondence, not a derivation. Most of the series' governance claims sit here: the mathematics is real; the mapping from mathematics to ministries is an argument, not a theorem.
- **[H] Heuristic** — estimated, illustrative, or applied loosely. Useful for orientation, not for adjudication. Index weights, the variety-ratio shorthand, and several cross-case analogies are explicitly

of this kind.

The distinction matters because the framework's most seductive failure mode is to let a rigorous formal result lend borrowed authority to a heuristic political reading. Naming the tier on each claim is the discipline that prevents this. Where a paper's formal interior is rigorous but its governance reading is in-principle, both labels appear.

The Thesis in One Move

Treat any governance system as a controller in a loop with the world it governs. It observes the world's state through a channel that selects some dimensions and drops others, decides on the basis of what it observes, acts, and observes the result. Three structural quantities then set hard limits, independent of competence or will: **latency** (the dead time between disturbance and response, which caps how fast the system can react), **signal fidelity** (how accurately information reaches the decision layer, which determines whether decisions track reality or a distorted image of it), and **dimensionality** (how many independent aspects of the world the system can perceive, which fixes the boundary between what can be governed and what arrives as a surprise). **[IP]**

From dimensionality comes the series' organising result, the **Goodhart–Ashby synthesis**. The claim is *not* that a controller must match its environment's complexity — it need not, and cannot: all regulation runs on reduced models, and a well-chosen low-order model is the normal and correct state of affairs. The claim is narrower, and it is about *which* disturbance dimensions a controller can perceive at all. Ashby's Law of Requisite Variety, in its disturbance-relative form, says a controller can only reject the variety of disturbances it can match; Goodhart's Law says a measure that becomes a target stops measuring what it tracked. Put together: the objective function selects which dimensions the system observes, and optimisation pressure then erodes its sensing of the dimensions it does not value — so the reduction is not a free engineering choice made from full sight, but one the controller cannot audit, because it cannot see what it has dropped. Those unvalued dimensions do not stop operating; they accumulate as externalities until they force a reckoning the system's own channels could not anticipate. The **variety gap**, G , is this objective-relative shortfall — the disturbance variety the architecture leaves unobserved — not the always-large, always-benign gap between the environment's raw complexity and the model. **[IP]** (*The static condition — that perceived variety must be at least the disturbance variety net of the objective's reach — is $[R]$ within the model; its identification with real institutional blindness is $[IP]$.*)

The consequence is that most reform changes the people, procedures, or resources *inside* an architecture without changing the architecture, and so leaves the ceiling in place. The failure is structural, and it is the structure the series tries to give a grammar.

What the Diagnosis Looks Like on a Case

Consider a health system trying to detect an emerging epidemic. Whether it sees clustered hospital admissions in real time or national mortality statistics months later is its *observation channel* and its *latency*. Whether a frontline clinician's report of an unusual case reaches the decision layer intact, or is averaged into a regional summary that dissolves the cluster into the mean, is its *signal fidelity* and its *representation chain*. Whether it models transmission in neighbouring jurisdictions or treats them as external shocks is its *boundary*. None of these is a question of competence or funding; each is a structural property the framework names, and together they decide what the system can and cannot see before anyone makes a decision. The grammar below is the catalogue of such properties.

Beneath the Grammar: The Observation Channel's Origin (Paper 0)

The thesis above, and every paper that follows it, begins after a step has already happened: the world's continuous flux has been carved into a finite set of variables — *these* count, *those* do not — and the controller reasons in *those*. That carving is the **factorization**, the observation channel itself, and the grammar below takes it as given. Paper 0, written late but meant to be read first, supplies the missing floor by asking where it comes from, and argues that factorization is not primitive but *derived*: two ingredients — bounded representational capacity and a bare temporal-prediction objective — suffice to force it, with no action, reward, or survival pressure required, so that a passive predictor already carves the world at joints that track its structure. A minimal model exhibits three registered properties **[R within the model]**: a bottlenecked predictor recovers latent causal variables it was never told about (*emergence*); under capacity starvation it does not blur evenly but sacrifices a coherent causal subspace whole, becoming a *static observer* that sees where things are but not where they are going (*structured blindness* — Paper VI's variety gap shown as a mechanism rather than a stipulation); and at the capacity margin it commits discretely to one of a competing pair of variables, the choice made by training contingency where the world is symmetric (*non-unique, symmetry-broken selection*). Two consequences ground what the later clusters build on: factorizations form large behaviorally-equivalent classes, so institutional coordination is selection *within* a class rather than discovery of a unique truth — yet within that freedom, classes that preserve the environment's causal variables are objectively better on robustness, sample efficiency, and transfer, so the world constrains the choice without fixing it. The two-ingredient argument is **[IP]**; the mechanism is **[R within the model]**, and its first registered operationalization of structured blindness failed and was corrected under a second preregistration, reported as part of the result.

The Grammar: Two Theory Cycles

Paper 0, written late but meant to be read first, sits beneath the whole structure as its foundational floor (tagged **Fdn.** in the table) and heads it. Of the twenty-seven papers that build on it, fifteen fall into two cycles. The first establishes the **static architecture** — the structural primitives a governance system has

whether or not it is changing. The second establishes the **dynamics of adaptation** — how an architecture changes, learns, and stays viable, or fails to. The table states each paper in the terms its own conclusion uses, with the tier of its governance claim. The two cycles now meet in a symmetry: the static deficits of the first compound multiplicatively (V), while the dynamic capacities of the second are gated by their minimum (XV) — the two ways a multi-part architecture's parts fail to be independent. The remaining twelve papers form two later clusters, described after the two cycles: a *limits-of-perception cluster* (XVI–XVIII), which interrogates the foundations the cycles build on, and a *factorization cluster* (XIX–XXVII, tagged **Factor** in the table), which works out the consequences of bounded representation. The evidential character of both differs from the cycles' and is discussed with each.

Paper	What it established	Cycle	Tier
0	Factorization — the observation channel every later paper takes as given — is not primitive but <i>derived</i> : bounded representational capacity plus a bare temporal-prediction objective suffice to force it, with no action, reward, or survival pressure. A bottlenecked predictor recovers latent causal variables it was never told about (emergence); under capacity starvation it sacrifices a coherent causal subspace whole rather than degrading evenly (structured blindness — the variety gap mechanized); at the margin it commits to one of a competing pair by training contingency, not the world (non-unique, symmetry-broken selection). Factorizations form large behaviorally-equivalent classes, so coordination is selection within a class, not discovery of a unique truth — yet classes preserving the environment's causal variables are objectively better.	Fdn.	R (within model) / IP
I	Latency and signal fidelity place hard ceilings on responsiveness; centralised controllers respond to the mean, not the distribution (the averaging problem).	One	IP
II	Disturbances arrive at many frequencies; a single-speed architecture is structurally mismatched to most of them.	One	IP
III	Representation chains attenuate the citizen-preference signal; beyond a critical depth the policy layer cannot reliably reconstruct it (constitutional unobservability).	One	IP
IV	Requisite variety lives at the point of contact; proximity is an engineering requirement before it is a political preference.	One	IP
V	The constraints of I–IV do not add — they compound multiplicatively (the coordination failure tax).	One	R / IP
VI	Objective functions are observation architectures; what a system does not value, it ceases to see (the variety gap).	One	IP
VII	Across fifteen national studies, reform disappoints structurally — through the institutional immune system, the bypass trap, and the legibility problem; the convergent first step is the protected experimental space.	One	H

Paper	What it established	Cycle	Tier
VIII	The variety gap made estimable: structural primitives, proxies, and a composite index G with stated uncertainty.	One	IP (index); H (weights)
IX	Architectural change as contested control: the latency asymmetry between reformers and incumbents, and <i>transition bandwidth</i> as a race that can be lost before it is visible.	Two	IP (Ω ratio: H)
X	Distributed sensing fails through correlation, not individual error: ensemble variance scales as $\sigma^2 \cdot [(1-\rho)/N + \rho]$. Its central prediction — that a population of contemporary AI observers is near-perfectly correlated — has now been preregistered and tested (Study 1, below).	Two	R (variance); empirically tested
XI	The actuation channel: delegation chains lose one dimension per deficient layer (a theorem), and control effort grows superlinearly with depth — yielding <i>constitutional uncontrollability</i> as the exact dual of III's unobservability.	Two	R (codimension); IP (energy)
XII	Boundary selection as an independent design variable; the small-gain theorem sets when unmodelled cross-boundary dynamics destabilise a well-run controller; the pooling paradox makes the trade-off inescapable for any single boundary.	One/Two	IP
XIII	Legitimacy as the first <i>endogenous coupling state</i> : it multiplies actuation and divides observation noise, cannot be set directly, and — when borrowed rather than built — collapses with hysteresis.	Two	IP
XIV	Stable learning as the third adaptation requirement: dual control, the exploration–exploitation tension, and persistent excitation as the rigorous content of "antifragility."	Two	IP
XV	The adaptation triad has finite throughput: sensing, learning, and execution share one recursive loop whose adaptive rate is gated by its slowest stage (the <i>adaptation bottleneck</i>), the dynamic dual of V's compounding; information, innovation, and reality backlogs are its failure signatures.	Two	R (bottleneck) / IP
XVI	Exploration-preservation resists formalization: across four disciplines a quantity representing unused alternatives decays under optimization and persists only through a <i>source term the optimizer does not set</i> . The order parameter is <i>source-term locality</i> — inside or outside the optimizer's control set — not the optimizer's "reach." A deliberately bounded, negative result.	Limits	IP
XVII	Processing can be made arbitrarily verifiable; <i>certification of reality</i> cannot be made self-verifying. Automating a coordination boundary relocates its irreducible world-certification link upstream but cannot delete it (the <i>relocation invariant</i>) — for world-coupled coordination, which governance is by construction; pure convention escapes it.	Limits	IP

Paper	What it established	Cycle	Tier
XVIII	Boundaries drift under the controller's own learning: when adaptation acts through channels that also carry cross-boundary coupling, <i>no fixed jurisdiction/environment decomposition survives learning</i> (the Non-Factorizability Theorem). A <i>Critical Learning Bandwidth</i> — too slow loses the plant, too fast dissolves the boundary — pinches shut on part of parameter space (the <i>Decomposability Frontier</i>). The registered early-warning index failed as defined and is reported as such.	Limits	R (within model) / IP
XIX	A held factorization carries three dissociable kinds of value — <i>govern</i> well, <i>warn</i> early (sentinel), or <i>bridge</i> the ecology — and governing skill predicts neither of the others; winner-take-all selection sheds warning and connective value as a side effect. Replicated across twenty retrained ecologies.	Factor.	R (within model) / IP
XX	Ashby, Goodhart, and certification cost are one bound seen under three operations — holding, optimizing, maintaining. The sharpened Goodhart: optimization erodes the target only when the projection discards a target-relevant dimension the optimizer can <i>reach</i> . The search for a conservation law is reported as failed.	Factor.	R (within model) / IP
XXI	The lifecycle of a bounded controller through three separations — learning $\neq$ adaptation, meta-learning $\neq$ free, persistence $\neq$ purpose — each a slide from an internal success to the external one it exists to serve.	Factor.	R (regress) / IP
XXII	Three assurances an institution wants — from inside, from outside, before the world — each fail, but from <i>three different sources</i> ; the "three limits from one bound" symmetry is refused. An institution cannot monitor its own certification kernel with instruments that depend on it.	Factor.	R (theorems) / IP
XXIII	The promised geometry of factorization space mostly does not exist (three registered failures): stress rescales it rather than reshaping it. What survives: behavioral distance is a symmetric metric, but <i>reform cost is directed</i> — asymmetric, non-composing, cheapest by an oblique route through the target's neighborhood.	Factor.	R (within model) / IP
XXIV	A diversity proxy tracks reach to valued options when observed passively but decouples once optimized: the optimized agent maxes the proxy while losing the reach. A persistent-excitation failure closed into an objective.	Factor.	R (within model) / IP
XXV	Bode's sensitivity integral imported as the conservation law behind "suppressed pressure relocates," with its limits marked: conserved log-area does not determine realized loss, which a strategic adversary reads off the upper tail. Exploitation needs concentration, accessibility, and discovery — none implied by conservation.	Factor.	R (Bode) / IP

Paper	What it established	Cycle	Tier
XXVI	Paper X's fixed escape-probability placeholder replaced by a competence state: escape from epistemic monoculture is <i>nucleated</i> by the best-preserved channel and <i>propagated</i> by a ladder cascade; institutional time enters as accumulated escape opportunity.	Factor.	R (within model) / IP
XXVII	Requisite variety is necessary but not sufficient. In an exactly-solved control problem, feedback that carries full information about the state but resolves a systematically <i>displaced</i> target loses substantial decision value (attenuation in 82% of tested conditions, no <i>enrichment</i>). Two displacements carrying identical information impose different costs, and the ordering reverses as displacement deepens — so <i>requisite alignment</i> (which distinction the signal resolves, not how many) is a second, geometrically structured constraint the Goodhart–Ashby synthesis did not name. The series' first computational-mechanistic paper.	Factor.	R / R (within model) / IP

A few results earn the **[R]** label outright. The multiplicative compounding of failure modes in V is arithmetic. The ensemble-variance result in X is a clean consequence of how correlated errors combine, and it is the formal heart of the claim that a single all-seeing model is more dangerous than many imperfect ones. The codimension law in XI — that a delegation chain loses exactly one cleanly transmitted dimension per deficient layer — is a theorem, not a numerical observation. And the adaptation bottleneck of XV — that a recursive loop's adaptive rate is the minimum of its stage rates — is likewise arithmetic, the dynamic dual of V's compounding. Most of the rest is **[IP]**: the control theory is sound, and the reading of institutions through it is a disciplined argument that a skeptical reader should treat as such.

The second cycle resolves into a single sequence, the **adaptation triad**:

Sense
↓
Learn
↓
Execute

Sense (Paper X) is keeping observers decorrelated enough to catch the error they share; *Learn* (Paper XIV) is exploring enough to keep the model identifiable; *Execute* (Paper IX) is retaining enough transition bandwidth to change the architecture in time. A system that cannot sense reality cannot learn; one that cannot learn cannot adapt; one that cannot execute its adaptation cannot survive. Sense → Learn → Execute is the series' answer to how a governance system stays adequate to an environment that generates novelty faster than architectures are usually redesigned.

The Reflexive Turn

Read together, the second-cycle papers converge on something the series did not set out to prove but kept rediscovering: a viable governance architecture must maintain the conditions of its own continued adaptation. Transition bandwidth (IX), observer diversity (X), legitimacy (XIII), and the capacity to learn (XIV) are not four unrelated requirements; they are four faces of one. Paper XIV names it directly — a controller that does not merely regulate the system but regulates its own regulation, the second-order move of cybernetics. This is offered as an emergent pattern, not a new primitive; adding it to the catalogue would be theory inflation. **[IP]**

It has one consequence worth stating plainly, because it falls out of the engineering rather than from any prior commitment. The activities that maintain a controller's model of a changing world — exploration, dissent, independent observation — are locally inefficient and globally necessary. Persistent excitation (XIV) is the formal statement that a system which suppresses all variance cannot identify its own parameters; observer decorrelation (X) is the formal statement that agreement among identical observers is not evidence. Protecting such activities is, on this reading, a design requirement for keeping the model calibrated, not a matter of taste. The framework makes the requirement visible; it does not pretend to derive from it any particular account of which ends governance should serve.

The Limits-of-Perception Cluster (XVI–XVIII)

Three later papers form a distinct group. Where the two cycles build the framework, this cluster interrogates the foundations the second cycle's adaptive capacities rest on: whether exploration can be preserved against the optimization that erodes it (XVI), whether the boundaries a controller trusts can be made durable against the optimizer that gains by moving them (XVII), and whether a jurisdiction stays separable from its environment once the controller's own learning reshapes the coupling (XVIII). Its character is deliberately different from the constructive papers — the findings are bounded, several are negative, and the highest confidence label is withheld by design. Each was produced by an adversarial multi-model protocol: a hypothesis routed through several frontier language models used as decorrelated observers and narrowed under sustained critique until only the claim that survived contact remained. Each records that its evidence is strong on "several disciplines formalise the same structure" and weak on "the structure is true of the world," because the observers are themselves correlated language models — the same caution Study 1 measured, now turned on the cluster's own method.

Paper XVI asks whether the many domain-specific warnings that success erodes adaptability name one object or merely rhyme. Four disciplines — control theory, evolutionary biology, institutional economics, decision theory — were each asked the same structural question in strictly native vocabulary, with the shared terms forbidden and the answer withheld. The independent formal statements share one skeleton: a quantity representing currently-unused alternatives decays under the primary objective and persists only through a *source term the optimising process does not itself set*. The corrected order parameter is **source-term locality**

— whether that term lies inside or outside the optimiser's control set. Two lenses place it strictly outside (a controller cannot enrich its own reference; an incumbent cannot cap its own barriers), two in a contested region the field has not closed. The paper explicitly refuses a unified theory, a measurable index, and a no-go theorem; the permitted claim is the shared decay-plus-source-term structure and the single axis along which the four differ. [IP]

Paper XVII isolates why such an external source term can be irreducible. The core is an asymmetry: **processing** — whether a system computes what it claims — can be made arbitrarily verifiable, but **certification of reality** — whether the external fact a rule depends on actually obtained — cannot be made self-verifying, because a verifier of a world-fact needs a verifier in turn, and the regress terminates only by trusting one anchor unverified. From this follows the **relocation invariant**: automating a coordination boundary relocates its irreducible world-certification link upstream but does not delete it — an immutable smart contract closes the execution link and reopens the same dependency at the specification link (what the tokens represent, what the proposal means). The invariant is scope-bounded: it binds *world-coupled* coordination, which must answer to some fact outside itself, and demonstrably fails for *self-referential* coordination (pure convention), where rule and practice are the same thing. Governance is world-coupled by construction, so the invariant applies to it. The design lever is not to eliminate trust but to minimise the chain to a single discrete certification link and cost-harden it — and to notice that hardening the *record* (a tamper-proof ledger) is not hardening the *certification* (the human or sensor attesting the recorded fact obtained, still the softest point). [IP]

Paper XVIII returns the cluster to an engineering register by removing an assumption from the boundary analysis of Paper XII. XII treated cross-boundary coupling as exogenous — a landscape the controller chases. XVIII treats the case where the controller's own learning reshapes the coupling, and proves the **Non-Factorizability Theorem [R within the model]**: under generic persistent learning, no time-invariant decomposition of the world into jurisdiction and environment survives — the boundary that was correct at design time is un-corrected by the adaptation it encloses. A minimal model exhibits the consequence as a cycle of factorizable calm, hidden coupling accumulation, non-factorizable collapse, and miscalibrated recovery, terminating past a reflexivity threshold in a *locked* state where the boundary welds shut. A **Critical Learning Bandwidth** bounds the learning rate from both sides — too slow loses the plant (Paper XV's bound in local form), too fast dissolves the boundary — and the two bounds pinch endogenously, closing entirely on a substantial region of parameter space: the **Decomposability Frontier**, a trade-off with an infeasible side. The reframing is that governance's boundary problem is not choosing the right boundary but maintaining a decomposability margin under the drift learning inevitably induces. [R within the model] / [IP]

XVIII also produced the series' clearest registered null. It committed in advance to an early-warning index — rising cross-boundary *prediction-error* correlation would flag boundary dissolution before it surfaced in outcomes — with the binding consequence that failure would cost the relevant sections their support. It failed: at a strict false-alarm budget the index as registered caught roughly one collapse in five. The

mechanism is the result. Local adaptation absorbs cross-boundary coupling into internal gain and thereby launders the evidence of dissolution out of exactly the residuals a well-run institution watches — Goodhart's law applied to the diagnostic itself, and a new instance of XVI's source-term locality now consuming the informativeness of a signal. The coupling survives only where no local objective is scrubbing it, in cross-boundary *state* covariance, and even there detection is lead-limited. The index was rewritten on state covariance and demoted to **[H]** pending field data. That a series about the limits of institutional perception found its own proposed instrument subject to them, and that the registered-prediction discipline caught it, is offered as the method working as intended.

The Factorization Cluster (XIX–XXVII): Bounded Representation and Its Consequences

A third group of nine papers works from a premise the earlier cycles left implicit and a foundational note (referred to across the group as *Paper 0*) makes explicit: a bounded controller cannot represent its world whole, so it must *factorize* — carve the world into a finite set of task-relevant distinctions — and the factorization it settles on is non-unique and regime-dependent, privileged by the environment rather than fixed by it. The cluster asks what follows once that premise is taken seriously: what a controller should do with the several factorizations it could hold, which of the series' borrowed laws are consequences of the bound rather than independent posits, how a bounded architecture ages, which assurances it can never give itself, what the geometry of factorization space does and does not support, and whether carrying sufficient variety is even enough for an observation to be worth acting on. Its character continues the limits cluster's discipline and sharpens it: almost every claim is carried by a preregistered simulation with committed thresholds and committed nulls, retrained across many seeds so that a result counts only if it survives the models being regenerated out from under it. The empirical claims are tagged **[R within the model]** — exact for the stated system, claiming nothing directly about any institution — with the governance readings held at **[IP]**. The cluster's most distinctive feature is the density of its honest nulls: several papers lead with predictions that failed, on the standing view that a registered falsification is a contribution, not an embarrassment.

Papers XIX and XXIII are an empirical pair on the portfolio of factorizations a controller holds. XIX asks what winner-take-all selection discards when it keeps only the currently best-governing model. Across twenty independently retrained model ecologies it finds that a factorization carries at least three dissociable kinds of value — it can *govern* well, *warn* early (a sentinel that sees the failure the active governor cannot), or *bridge* the ecology (hold open translation paths between factorizations that would otherwise drift out of reach) — and that governing skill predicts neither of the other two. An adaptive audit that preserves all factorizations for sensing but commits to one for action, reopening on evidence, tracks an unattainable oracle and beats a monoculture in all twenty ecologies; the top governor differs from the top bridge in all twenty and from the top sentinel in seventeen (though the governor–sentinel correlation, at $\rho = 0.48$, makes that dissociation real but weak). Two registered claims failed and are reported as failures, disciplining the pilot's

overclaims. The result gives XVI's source terms a mechanism — optimization sheds early-warning capacity not because warning is costly but because it is uncorrelated with governing skill — and adds a failure mode the earlier papers lacked: governance can fail by fragmentation of factorization space, which no improvement of the active governor repairs. XXIII, the promised geometric sibling, opens by reporting that the object it was to describe mostly does not exist: a pre-drafting falsifiability audit and registered replication found that stress *rescales* factorization space rather than reshaping it, that the connectivity threshold XIX had leaned on is provably just a restatement of distance magnitude, and that no topological transition appears under a continuous stress sweep. What survives is sharper for the failures: behavioral distance is a symmetric metric, but *reform cost is directed* — strongly asymmetric (what an architecture costs to leave is not what it costs to return to), non-composing, and only weakly predicted by behavioral distance. Routing a reform through an intermediate lowers its cost, but a registered control shows this is not geodesic: the best intermediate depends on the reform's destination, not its origin — reform is cheapest by an oblique approach through the target's neighborhood. **[R within the model] / [IP]**

Papers XX and XXII are a formal pair on what the bound implies and, pointedly, what it does not. XX shows that three of the laws the series has leaned on — Ashby's requisite variety, Goodhart's law, and XVII's certification cost — are not independent posits but one bound seen under three operations: holding a factorization (Ashby, a pigeonhole theorem, nearly definitional and flagged as shallow), optimizing through it (Goodhart), and maintaining its fit over time (certification cost, a monotone accounting inequality). The one substantive new result is a sharpened Goodhart: optimization degrades the target *precisely when* the lossy projection discards a target-relevant dimension the optimizer can *reach* — reachability of the discarded dimension is the discriminator, which says which proxies are safe and which are not (demonstrated across thirty registered worlds, **[R within the model]**). The paper also reports a failed search for a *conservation* law: representational complexity can rise or fall; only the cost of staying aligned is monotone, and that negative result is stated rather than smoothed. XXII refuses the tempting parallel — "three limits from one bound" — and shows it is false. Three assurances an institution might want fail from three different sources: reform-convergence undecidability requires computational *universality* (the negation of bounded representation, since a finite controller has a decidable convergence problem); the No Free Lunch limit requires only the absence of a prior over environments and binds bounded and unbounded controllers alike; only certification incompleteness traces to the bound. In this terrain, XXII observes, rigour and interest run in opposite directions — the two clean theorems are shallow, the one result with real content is not a theorem. Its empirical demonstration failed on all four registered predictions, and what replaced them (re-registered, confirmed on twenty fresh seeds) is the cluster's sharpest practical claim: inverting an institution's need-detection channel does not starve the misidentified party but *floods* them, so unmet need reads a perfect zero and the failure is invisible from inside — an institution cannot monitor its own certification kernel with instruments that depend on that kernel, and a failing institution's health indicators can *improve*. **[R within the model] / [IP]**

Paper XXI treats the lifecycle of a bounded architecture through three separations, each a place where a controller mistakes an internal success for the external one it exists to serve. *Learning is not adaptation*: improving model fidelity and maintaining coupling to the world can oppose, and a registered minimal model shows the fastest learner holding the best model and the worst grip — coupling peaks at an intermediate learning rate and then degrades, gated by an absorptive-capacity inequality (**[R within the model]** for the separation). *Meta-learning is not free*: the ladder of learning-to-learn cannot regress indefinitely in a bounded system, so it must close on a small set of invariants held still — identity boundary, certification kernel, memory, timescale separation, a plural reserve — and constitutional engineering is, on this reading, the choice of what to hold fixed. *Persistence is not purpose*: an institution's terminal adaptive act may be to end or transfer its coupling rather than continue, and the characteristic failure is the controller that preserves itself in place of the coupling it was built to maintain. The last two are argued, not demonstrated, and marked as such. **[R] (regress) / [IP]**

Paper XXIV isolates a sharper form of Goodhart-as-sensor-corruption in a minimal adaptive system. A diversity proxy built on a coarse-grained observer tracks an agent's genuine reach to valued peripheral options *when exploration varies passively*, but decouples entirely once the proxy becomes an optimization target: the optimized agent drives the proxy to its ceiling while retaining no ability to reach the periphery. Peripheral reach is preserved only when the observer resolves the periphery finely enough relative to its access cost, with resolution and cost acting additively. And the bias need not be imposed — a representation learned from task-centred experience spends its resolution on the frequently visited centre and compresses the periphery on its own, landing in the decoupling regime; broader exploration does not repair this, but structural exposure that samples the periphery does. The mechanism is a persistent-excitation failure closed into an objective: optimizing the observer suppresses the excitation that would let it see what it is meant to measure. The setting is a single deterministic gridworld with no strategic agents; the contribution is a mechanism and a falsifiable prediction for administrative measurement, stated as hypothesis, not a law of governance. **[R within the model] / [IP]**

Paper XXV gives the series' recurring "suppressed pressure relocates" claim the formal backing it usually lacks, and is scrupulous about where the backing stops. It imports Bode's sensitivity integral — a genuine conservation law: for a stable loop of relative degree at least two, log-sensitivity conserves, so suppression in a monitored band lower-bounds amplification elsewhere (**[R]**, within the linear, fixed-controller, non-strategic conditions that do not literally hold in governance). Its conceptual content is a mismatch of geometry: conservation acts on the positive part of $\log|S|$, while a strategic adversary's realized loss acts on the upper tail of $|S|^2$ (the CVaR of the response over a reachable set). Because these are different functionals of the same profile, conserved pressure does not determine realized harm, and the accessible amplified *area* does not determine the accessible *loss*. Exploitation is licensed by three separable capacities — concentration, accessibility, discovery — none implied by conservation. A convex-synthesis study then shows that buying deeper monitored-band suppression raises the exploitability floor, and that a looser

complementary-sensitivity allowance lowers it only until suppression is maximal, where design freedom collapses. The paper is an explicit exercise in *not* committing the series' named borrowed-authority error: the conservation law is imported for exactly what it licenses and no more. **[R] / [IP]**

Paper XXVI removes a placeholder from Paper X. X represented the difficulty of escaping an epistemic monoculture by a single fixed near-zero probability of reverting to independence; XXVI replaces it with an explicit competence state that rebuilds under independent operation and decays under shared-system use. A homogeneous mean-field reduction solves in closed form, decomposing the consolidation–return hysteresis into a consensus term, a competence-decay term, and a liability term — with competence decay provably moving the *return* threshold but not the *entry* threshold. The heterogeneous population departs from that reduction in a structured way that is the paper's main object: escape from full consolidation is *nucleated* by the best-preserved channel and *propagated* by a cascade whose feasibility is a ladder condition on the ordered penalties (the k-th cheapest channel must be viable after k–1 defections). The resulting escape-ladder theorem yields a finite-population quantile criterion — no fitted phase boundary — that classifies fifty-five of fifty-six cells of a frozen grid, its single miss on the boundary itself, and a stochastic layer shows escape hazards composing into the observed exit thresholds so that institutional time enters as accumulated escape opportunity. Several registered predictions failed along the way and were replaced; the decomposition and ladder theorem are exact within the model, the transfer to real epistemic infrastructure **[IP]** at best. **[R within the model] / [IP]**

Paper XXVII asks whether dimensional sufficiency — requisite variety — is enough for an observation to be worth acting on, and finds it is not. Working in a single exactly-solved partially-observed control problem, it measures by preregistered dynamic programming the decision value of a scalar feedback signal — how far the controller's expected loss falls because it can pay to observe rather than act blind — and then *displaces* that signal so that it evaluates a systematically wrong target while remaining fully informative and correctly modelled by the optimal controller. Displaced feedback loses substantial value: resolved, substantive attenuation in 486 of 592 tested conditions (82%), unanimous among all conditions whose numerics resolved, and provably not reducible to a loss of channel reliability. A registered possibility that displaced feedback might *exceed* matched feedback — *enrichment* — was permitted by the design and did not occur. The decisive result is geometric. Two displacement geometries carrying identical marginal information about the state impose different value costs, and the ordering reverses as displacement deepens: one geometry preserves more value at low displacement, the other at high. Because information content is held exactly equal across the two, the difference cannot be an information-quantity effect; it isolates a constraint orthogonal to variety. The paper names it **requisite alignment**: for an observable to have decision value it is not enough that it carry sufficient variety and transmit it without compression — the distinction it resolves must line up with the distinction the controller's action turns on, and alignment has geometry (which distinction, not how many). The worst case is not maximal displacement but intermediate *ambiguity*, because a consistently wrong signal can be partially inverted where an unpredictable one cannot. This is the series' first computational-mechanistic paper, and its result refines the organizing Goodhart–Ashby synthesis

directly: requisite variety is necessary but not sufficient, and requisite alignment is a second, independently binding, geometrically structured constraint. The finding holds within one exactly-analysable case and estimates its prevalence in no governance system. **[R] / [R within the model] / [IP]**

The Third Cycle: From Theory to Test, and Then to Build

The first two cycles are theory. They diagnose, formalise, and measure in prototype. Between that theory and any engineering sits an **empirical gate**: a framework that declines to be confronted with data does not yet deserve the name engineering. The series' standing rule, since the measurement paper, is that a documented null is more valuable than an untested elaboration.

The first prediction has now been through the gate. Paper X predicts that contemporary AI systems, increasingly used as governance observers, are near-perfectly correlated — that consulting more of them buys almost none of the error reduction that independent observers would. Study 1 tested this under a frozen, preregistered protocol (battery published before collection; a blind external critique adjudicated; nulls committed to in advance). Six consumer AI systems each estimated fifty governance-relevant quantities sampled from public databases, scored against ground truth. The effective error correlation was $\rho_{\text{eff}} \approx 0.97$ (95% CI roughly [0.95, 0.99]): a six-model ensemble sat almost exactly at the single-model error level, where independence would have cut it roughly sixfold. The primary prediction held, decisively. The secondary prediction — that correlation would be strongest in the tails — was **not** supported, and is reported as such. The limits are stated in the protocol and matter: six systems, consumer interfaces rather than controlled instruments, items restricted to quantities predating the models' training cutoff, and the protocol designer among the subjects (disclosed, and mitigated by moving item selection to seeded draws from public databases). An ecological complement using recent quantities remains on the roadmap. **The claim this licenses is narrow and strong: the correlation-tax mechanism is real for current AI observers — not that the framework as a whole is validated.**

That is one prediction. The rest of the empirical programme is specified and open: a variety-gap pilot audit of a willing institution; the delegation-depth versus implementation-fidelity study at proper sample size; a prospective variety-gap panel across twenty to thirty governance systems; the legitimacy-estimation protocol applied to a representative sample. Beyond the empirical phase lies the engineering proper — the protected experimental spaces, the independent observer ensembles, the legitimacy sensors and circuit-breakers that the design principles point to. None of that has been built.

It is set down here as an **open invitation**, and for a reason consistent with the framework's own logic. The series argues that no single integrator should be the bottleneck on a system meant to perceive more than any one vantage can; Paper X makes the point formally, and a project developed through one editorial judgement is a candid instance of exactly that limitation. The comparative advantage of this effort has been diagnostic and formal. The build-out — pilots, instruments, institutional design, sustained empirical testing — is a different discipline, and it is open for collaborators whose strengths lie there to take up, modify, and improve.

The primitives are defined precisely enough to be operationalised; the diagnostic diagram can be completed for any new case as an exercise; the protocols, the simulations, and Study 1's frozen analysis script are reproducible.

What the Framework Does Not Claim

The honest boundary matters as much as the claims.

It does not supply the *content* of good governance. It specifies structural constraints on viable governance architectures; it does not determine the ethical ends those architectures should pursue. Following Habermas' distinction, it speaks to the facticity of institutions, not their validity. Engineering can design a viable vessel; it cannot decide where the vessel should sail.

It is not a general theory of systems. The mid-twentieth-century attempts at one — von Bertalanffy's general systems theory and its successors — promised more than they delivered, and the critique of that overreach is well taken (Berlinski, *On Systems Analysis*, MIT Press, 1976). The discipline taken from that history is deliberately modest: control-theoretic concepts earn their keep as diagnostic lenses on specific problems, in combination with deep domain knowledge, not as a unified theory from which governance outcomes can be deduced. Where this brief uses the word *grammar*, it means a catalogue of structural properties to check case by case — not a generative system that explains institutions on its own.

Nor is the core observation — that a single low-dimensional metric distorts — itself new: Goodhart's Law, Ridgway's "what gets measured gets done", and economics' decades-long turn toward multidimensional measures of development and welfare all say as much. The contribution offered here is the structural mechanism — the objective function as observation architecture, and its erosion under optimisation — not the observation that metrics distort.

Several quantities that read as precise are not. The composite index G is structurally motivated **[IP]**, but its tier weights and its critical threshold are **[H]** — parameterisations calibrated against the case set, not derived from first principles, and reported with sensitivity analysis for that reason. The variety-ratio shorthand of the transition paper is a heuristic and is kept out of public-facing material. The translation of the actuation paper's energy law into "political capital" is in-principle, never rigorous, and its falsifiable predictions are deliberately stated in fidelity and depth, which can be coded, rather than in energy, which cannot. Some axes — notably where a real delegation chain sits between its idealised poles — lack a field instrument entirely.

And the provenance is what it is. This is a recent, solo, AI-assisted project, not a long-standing research programme. The cross-case patterns are consistent enough across radically different domains that they are unlikely to be noise, and the formal cores are checkable. Study 1's finding bears on the project's own method as much as on anyone else's: the series was assembled with help from several of the same systems the study measured, so their agreement cannot be read as corroboration. What the process relied on was not their averaged estimates but the disagreements it could surface and an editor's integration of them — and that

reliance is a claim to be checked, not a defence to be assumed. The corpus is a starting point built by an architecture with one editor at its centre, and its most useful future is to be extended by people who can occupy positions its author cannot.

Invitation

The framework is a living diagnostic instrument, not a proprietary method. The useful responses to it are to test it against cases it has not seen, to challenge the primitives where they do not fit, to extend the formal foundations where they are thin, and — in the third cycle that has barely begun — to measure and to build. One prediction has survived contact with data. The work of building remains. The architecture for building is, at least, specified well enough to start.