




The Coordination Failure Tax

Architectural Compounding and the Path to Requisite Governance

Paper V in the Governance as Engineering series

Four governance failure modes — spatial blindness, frequency gaps, preference invisibility, and observational inadequacy — do not add. They multiply. A governance system exhibiting all four simultaneously is categorically incapable of the functions it claims to perform.

Björn Kenneth Holmström

March 2026

Creative Commons Attribution-ShareAlike 4.0 International

<https://bjorkennethholmstrom.org/working-papers/coordination-failure-tax>

Abstract

Papers I–IV establish four distinct structural failure modes in governance architecture: spatial blindness from aggregation (Paper I), frequency gaps from single-scale control (Paper II), preference invisibility from deep representation chains (Paper III), and observational inadequacy from low-dimensional commons monitoring (Paper IV). Each paper demonstrates that the failure mode it identifies is structural — immune to institutional improvement because it operates at the level of the observation channel, upstream of any signal that institutional quality can act on.

This paper makes explicit what the series implies but never states: these failure modes do not add — they multiply. A governance system exhibiting all four simultaneously is not four times worse than a well-designed one. It is categorically incapable of the functions it claims to perform. We introduce the *coordination failure tax* as a formal concept: the hidden cost imposed on any system operating below requisite variety across multiple architectural dimensions simultaneously. We show that this cost compounds multiplicatively, that existing empirical data from democratic preference studies and commons collapse records is consistent with the compounding prediction, and that each design principle in the Global Governance Frameworks corresponds to a structural response to one of the diagnosed constraints. The paper concludes by identifying which categories of institutional reform are architecturally capable of reducing the tax, and which are not.

I. What the series established

The four papers in the Governance as Engineering series address different governance domains — crisis response, multi-scale disturbance management, democratic representation, and commons governance — using different formal tools: control theory, frequency analysis, information theory, and Ashby's Law of Requisite Variety. Their surface diversity obscures a structural unity.

Each paper demonstrates the same underlying mechanism in a different domain. Aggregation destroys information. The destroyed information cannot be recovered downstream. Institutional quality operates on the signal after it arrives and cannot help if the signal is gone.

The four failure modes, seen as instances of this mechanism:

Paper	Domain	What is aggregated	What is lost	Result
I	Crisis response	Local conditions	Spatial information	Uniform response to diverse problems
II	Disturbance management	Frequency bands	Temporal resolution	Missing fast and slow dynamics
III	Democratic representation	Citizen preferences	Preference distribution	Policy governed by phantom signal
IV	Commons governance	Ecological dimensions	Multi-band resource state	Collapse despite monitoring

In each case, the observation channel between the governed system and the governing layer destroys the variety the governing layer would need to perform its function. Paper I shows this as a stability ceiling: above a certain latency-to-response-speed ratio, no controller gain can stabilize the system. Paper II shows it as a frequency gap: no single-scale controller can cover the full disturbance spectrum of a complex environment. Paper III shows it as a constitutionality threshold: representation chains deeper than two or three layers cannot transmit citizen preferences to the policy level because noise variance exceeds surviving signal variance. Paper IV shows it as a variety mismatch: governance systems with fewer observation dimensions than their resource environment has disturbance bands will systematically authorize extraction above sustainable yield.

The consistent finding across four domains, four formal frameworks, and four simulation architectures suggests the framework is capturing something real about governance — not a collection of domain-specific observations but a unified structural constraint on what governance can achieve under specified architectural conditions. The simulation parameters underlying each finding are open for inspection and challenge; the claims are falsifiable by changing the parameters and testing whether the results hold.

II. The interaction effect — why failures compound

The four failure modes are not independent. When a governance system exhibits multiple structural failures simultaneously, the failures amplify each other. The effective governance capacity of the system is not reduced by the sum of the individual deficits. It is reduced by their product.

The compounding mechanism

Before stating the relation formally, it helps to see why it must be multiplicative rather than additive. Each failure mode does not operate on a fixed baseline in parallel with the others — it operates on the output of the governance processes that precede it in the causal chain. Spatial blindness corrupts the signal before frequency management can act on it. Frequency gaps leave disturbance bands uncovered, generating unresolved variance that then contaminates the preference signal. Preference invisibility means that even an intact fractal architecture receives corrupted mandates at each level. Observational inadequacy means that even if preferences are intact, the resource state they would constrain is misread. Each failure mode takes as its input the already-degraded output of the others. The losses are sequential, not parallel.

G_0 — effective governance capacity at baseline — can be operationalized differently across domains: mean time to recovery from standardized disturbances in crisis management, preference-policy correlation under controlled conditions in democratic governance, or resource stock stability under standardized extraction pressure in commons management. What matters for the compounding argument is not the specific metric but that G_0 represents the capacity that would be available if all four observation channels were intact. The failure modes reduce what arrives through those channels.

Let f_1, f_2, f_3, f_4 represent the fractional capacity destroyed by each failure mode — the proportion of effective governance capacity lost to spatial blindness, frequency limitation, preference invisibility, and observational inadequacy respectively. The intuitive model treats these as additive losses:

$$G = G_0 - c_1 f_1 - c_2 f_2 - c_3 f_3 - c_4 f_4$$

This model is wrong. The failures are not independent subtractions from a fixed baseline. Each failure mode corrupts the input available to the governance mechanisms that the other failure modes leave intact. The correct relation is multiplicative:

$$G = G_0 \cdot (1 - f_1) \cdot (1 - f_2) \cdot (1 - f_3) \cdot (1 - f_4)$$

Each factor $(1 - f_i)$ is bounded between 0 and 1. When any f_i approaches 1 — when a single failure mode destroys nearly all governance capacity in its dimension — G approaches 0 regardless of performance in the other three dimensions. This is the single-point-of-failure property: total observational collapse in any one dimension eliminates effective governance capacity entirely, independent of how well the other channels are functioning.

When any single failure is present at $f = 0.5$, effective capacity is halved. When all four are present at $f = 0.5$, effective capacity is not halved but reduced to 6.25% of baseline — less than one-fifteenth of what the additive model would predict.

This is the coordination failure tax: the compounding cost of operating below requisite variety across multiple architectural dimensions simultaneously. Like a financial tax, it is extracted invisibly, continuously, and regardless of how hard the system works within its existing structure.

How each failure amplifies the others

The multiplicative relation follows from the causal dependencies between the four failure modes.




Figure 1. Information flow through a governance system: four failure modes positioned at their points of attack on the observation channel.

Spatial blindness amplifies frequency gaps. Paper I establishes that central governance systems receive an aggregated signal that destroys spatial information — the shock at node 3 is invisible to a controller that observes only the system mean. Paper II establishes that single-scale governance systems cannot cover the full disturbance frequency spectrum. When both failures co-occur, the frequency gaps are not merely uncovered — they are filled with noise. The aggregated signal carries distortions from unresolved local variation that appear at the central level as apparent disturbances at frequencies the central controller was not

designed to address. The result is not only coverage gaps but active interference between the spatial and frequency failure modes. Systems with both failures will experience not just delayed response to crises but wrongly targeted delayed response — interventions calibrated to the wrong place and the wrong timescale simultaneously.

Preference invisibility corrupts the inputs to fractal architecture. Paper III establishes that deep representation chains cannot transmit citizen preferences to the policy level. Paper II establishes that fractal nested governance is the stability-optimal architecture for multi-scale disturbance management. When both failures co-occur, the fractal architecture — even if present — receives corrupted inputs at each level. The local council that should respond to fast local disturbances receives a mandate derived from a preference signal that no longer resembles what local citizens actually want. The structural solution to the frequency gap problem requires intact preference transmission to function; preference invisibility undermines it from the inside. The practical consequence: a system with this pair of failures will exhibit the formal structure of distributed governance while functionally operating as a centralized one, because the local variation in mandate that distributed architecture requires is destroyed before it can guide local action.

Observational inadequacy and preference invisibility produce synchronized collapse. Paper IV establishes that commons governance systems with insufficient observation dimensionality systematically authorize extraction above sustainable yield, accelerating collapse. Paper III establishes that deep representation chains lose the citizen preference signal that would constrain extraction. When both failures co-occur, the governance system is authorizing unsustainable extraction at precisely the moment when the mechanism that would correct it — citizen preferences for conservation reaching the policy level — has been destroyed in transmission. The two failure modes synchronize their effects: one removes the brake, the other removes the warning that the brake is gone. When extraction is authorized by a resource-blind system and constrained only by a preference-deaf one, the remaining bound on extraction is physical collapse.

Empirical consistency

These interaction effects are predictions, not observations. But existing empirical data is consistent with them in ways that the individual failure modes alone would not predict.

Gilens and Page's 2014 analysis of 1,779 policy issues in the United States found that average citizen preferences have near-zero statistical influence on policy outcomes, while economic elites and organized interest groups retain substantial influence. The magnitude of this divergence is what Paper III's SNR threshold would predict for a representation system of the observed depth operating with realistic noise parameters. It is not, by itself, proof that the SNR mechanism is operating — alternative explanations exist, critics argue that the interest group variables are noisy and that results are sensitive to model specification, and the study is confined to the United States. But the pattern is consistent with the compounding prediction: a system exhibiting both deep representation chains (Paper III) and governance concentrated in actors with privileged observation access (Paper I) would be expected to show exactly this preference-policy divergence, for structural rather than motivational reasons.

The FAO's global fisheries data presents a complementary pattern. Of assessed fish stocks, approximately one-third are currently overfished and over half are fished at maximum sustainable levels. State management — the predominant governance architecture for marine commons — is the Architecture B that Paper IV's simulation identifies as performing worse than open access under certain observation latency conditions. The mechanism the paper identifies (annual aggregate quota calibrated to last year's stock, authorizing overharvest in declining resources) is documented in fisheries management literature as a persistent structural failure, not a failure of enforcement or institutional commitment. The North Atlantic cod collapse is the paradigm case: annual stock assessments and scientifically-derived quota recommendations failed to prevent collapse because the observation lag between stock decline and quota adjustment allowed overharvest to accumulate to the point of population collapse — exactly the high-latency single-dimension failure mode Paper IV models. The data is consistent with the Paper IV prediction that single-dimension, high-latency observation produces characteristic collapse dynamics regardless of institutional quality.

Neither dataset constitutes proof of the compounding mechanism. Both are consistent with what the compounding mechanism would predict, in the domains where that mechanism is most clearly specified. The appropriate frame is not "this data confirms the theory" but "if the theory is correct, this is what we would expect to see, and this is what we see."

What would constitute stronger grounding — empirical measurement of actual latency distributions, aggregation ratios, and noise levels in specific governance systems — remains future work. The simulation methodology established across Papers I–IV provides a clear protocol for that measurement: the parameters the simulations use are estimable from administrative data, and their relationship to governance outcomes is now formally specified.

The compounding mechanism established in this section is a diagnostic instrument. It identifies not only that existing governance systems underperform but why they underperform and by how much — and it makes a further prediction: that systems exhibiting all four failures simultaneously will perform qualitatively worse than the sum of their individual failures would suggest. Section III examines what an architecture designed to satisfy all four structural constraints simultaneously looks like, and shows that each design principle in such an architecture is load-bearing against a specific diagnosed failure mode.

III. The GGF as a structural response

The Global Governance Frameworks constitute a proposed post-Westphalian coordination architecture. They are often described in normative terms — as a response to climate breakdown, inequality, and institutional failure. This framing is accurate but incomplete. The GGF's design principles are not primarily normative commitments. They are engineering responses to the four structural constraints the series has diagnosed. Each load-bearing element of the architecture addresses a specific failure mode. Remove any one element and the corresponding failure mode goes unaddressed; the coordination failure tax for that dimension is not reduced.

This section makes the mapping explicit.

The Westphalian baseline

Before examining the GGF's responses, it is useful to characterize what they are responding to. The Westphalian nation-state system — the dominant governance architecture for most of the world — exhibits all four failure modes simultaneously.

Observation is concentrated at the national level, which receives aggregated statistics about local conditions. A housing shortage in one district, soil degradation in one watershed, or a community resilience deficit in one region is invisible to national governance until it has grown large enough to appear in national aggregates — by which point the problem has already compounded. This is the Paper I failure mode: spatial blindness from aggregation.

National governance operates at a single temporal scale, calibrated to electoral cycles and legislative processes. Fast-moving crises — pandemic emergence, financial contagion, local ecological tipping points — outrun the national response timescale. Slow-moving structural changes — soil degradation, demographic transition, long-run deindustrialization — move below the threshold of electoral visibility. Neither the fast nor the slow frequency band is adequately governed. This is the Paper II failure mode: frequency gaps from single-scale control.

Democratic representation in Westphalian systems typically operates through chains of three to five layers: citizen to local representative to regional assembly to national parliament to executive. Paper III establishes that chains of this depth cannot transmit citizen preferences to the policy level with positive SNR. The policy layer governs a phantom signal — a distorted echo of citizen preferences that has been aggregated, averaged, and noise-corrupted beyond recovery. This is the Paper III failure mode: preference invisibility from deep representation chains.

Natural commons — fisheries, forests, watersheds, atmospheric sinks — are predominantly governed by national and international bodies using aggregate monitoring and periodic quota-setting. Paper IV establishes that this architecture performs worse than open access under conditions of resource decline, because the

observation lag between stock change and quota adjustment authorizes overharvest precisely when the resource can least afford it. This is the Paper IV failure mode: observational inadequacy from low-dimensional commons monitoring.

The Westphalian system, in other words, is not a system that partially satisfies the structural constraints. It exhibits all four failure modes simultaneously. The coordination failure tax is not reduced at any dimension. The compounding calculation from Section II applies in full.

The GGF's structural responses

Figure 2. The GGF as a structural response to four diagnosed failure modes
Each architectural principle is load-bearing against a specific constraint

Paper / Failure Mode	Westphalian Manifestation	GGF Structural Response
Paper I Spatial blindness from aggregation	National aggregate stats; local crises invisible until nationally visible	Subsidiarity + Bioregional Governance Units Governance embedded within ecological boundaries; observation latency approaches zero for local disturbances
Paper II Frequency gaps from single-scale control	Electoral cycles miss fast crises and slow structural change simultaneously	Polycentric Nested Architecture Community → bioregional → continental → global layers, each calibrated to its frequency band
Paper III Preference invisibility from deep chains	3–5 layer representation; SNR below threshold; policy governs phantom signal	Direct Participation + Shallow Chains Community assemblies feed councils through ≤2 layers; Circuit Breakers validate preference signal at each level
Paper IV Observational inadequacy (commons)	Annual aggregate quota; single-dimension signal; authorizes overharvest	Indigenous Wisdom Council + Commons Trusts Multi-generational slow-variable signal; continuous multi-dimensional monitoring within resource boundaries

Figure 2. The GGF as a structural response to four diagnosed failure modes.

The GGF addresses each failure mode through a specific architectural choice. The mapping is not incidental — each design principle is the structural response to the diagnosed constraint.

Paper I failure (spatial blindness) → Subsidiarity and bioregional governance units

The GGF places governance authority within the systems it governs. Bioregional Autonomous Zones are defined by ecological boundaries — watersheds, biomes, coastal systems — rather than administrative ones, and are governed by councils embedded within those boundaries. This positioning eliminates the aggregation step that destroys spatial information in Westphalian governance. The council observing a watershed is positioned within the watershed; it observes conditions directly rather than through an aggregate signal. The observation latency approaches zero for local disturbances; the spatial information that would be destroyed in aggregation is never aggregated.

The subsidiarity principle — that decisions should be made at the lowest level capable of addressing the problem — follows from the same structural analysis. It is not a preference for local governance as such. It is the recognition that governance positioned outside a system cannot observe what is happening inside it, and that information destroyed in transmission cannot be recovered at the receiving end.

Paper II failure (frequency gaps) → Polycentric nested architecture

The GGF's polycentric structure — overlapping jurisdictions of Territorial Councils, Commons Trusts, Guilds, and the Indigenous Wisdom Council, operating at community, bioregional, continental, and global scales — is designed to cover the full disturbance frequency spectrum that no single governance scale can cover alone.

Fast disturbances — acute local crises, rapid ecological change, community-level conflict — are addressed by community-level governance with short response latency. Medium disturbances — seasonal resource variation, regional economic cycles, multi-year ecological trends — are addressed by bioregional governance calibrated to those timescales. Slow disturbances — multi-decadal ecological shifts, civilizational transition pressures, intergenerational resource dynamics — are addressed by the Indigenous Wisdom Council and global coordination layers operating at the timescales at which those disturbances manifest.

The architecture does not produce redundancy across scales. It produces coverage — each scale governing the frequency band it can actually observe and respond to, without attempting to govern the bands that its observation latency makes inaccessible.

Paper III failure (preference invisibility) → Direct participation and shallow chains

The GGF's democratic mechanisms are designed around the SNR constraint established in Paper III. Where Westphalian systems route preferences through chains of five or more layers, the GGF's participatory structures minimize chain depth: direct community assemblies feed bioregional councils through a single layer of aggregation, not through multiple layers of representation. The preference signal survives this transmission because the chain is short enough to keep noise variance below signal variance.

Digital participation infrastructure — where implemented — further reduces effective chain depth by enabling citizens to express preferences on specific decisions rather than delegating their entire preference set to a representative for a fixed term. The preference signal becomes more specific, more frequent, and less aggregated, reducing the information loss at each step.

The Circuit Breaker mechanisms — which pause implementation of decisions that exceed specified thresholds of community dissent — function as preference signal validation: they detect when the policy layer has drifted beyond the boundary of what the preference signal, however imperfectly transmitted, would authorize.

Paper IV failure (observational inadequacy) → Indigenous Wisdom Council and Commons Trusts

The GGF's most distinctive architectural feature, from the perspective of the series, is the structural role assigned to indigenous governance systems. Paper IV establishes that effective commons governance requires observation dimensionality matching the resource system's disturbance variety — and that the slow ecological signal, which spans decades to centuries, can only be accessed by governance systems that have been embedded within specific ecologies across those timescales. Indigenous governance systems, developed through generations of direct resource observation, possess exactly this slow-variable observation capacity. The knowledge is not metaphorical or cultural. It is a multi-generational observational record of slow-moving ecological variables that modern monitoring programmes have not existed long enough to observe.

The Indigenous Wisdom Council is not, in this framing, a concession to identity politics or historical redress — though it may be those things as well. It is a structural requirement for commons governance with sufficient observation dimensionality to cover the slow frequency band. An architecture without it has a permanent frequency gap at the slow end of the disturbance spectrum, regardless of how sophisticated its other monitoring systems are.

Commons Trusts — governance bodies holding renewable resources in stewardship rather than ownership — address the same failure mode at the medium frequency band. By positioning governance authority within specific resource systems with continuous monitoring mandates, they provide the multi-dimensional observation that periodic national survey cannot.

What happens if any element is removed

The load-bearing nature of each architectural choice becomes visible when any one element is removed.

A GGF without the Indigenous Wisdom Council — retaining subsidiarity, polycentric structure, and shallow representation chains — still fails on observation dimensionality. The slow ecological signal remains inaccessible. The Paper IV failure mode is unaddressed. For any commons governed by this truncated architecture, the simulation predicts collapse risk approaching Architecture B levels for disturbances at the slow frequency band, regardless of how well the other three constraints are satisfied.

A GGF without shallow representation chains — retaining bioregional positioning, polycentric frequency coverage, and commons observability — still fails on preference transmission. The policy layer governing each bioregional council is separated from citizen preferences by a chain deep enough to lose signal. The community assemblies and commons trusts operate with high observation dimensionality; the preference signal that should guide their mandate is still being destroyed in transmission. The structure is well-observed but ungoverned by actual citizen preferences.

Each element is necessary. None is sufficient alone. The architecture is integrated in the specific sense that satisfying three of the four constraints while leaving the fourth unaddressed still pays the coordination failure tax on the unaddressed dimension.

IV. Transition pathways

The diagnosis in Sections I–II and the design specification in Section III together create a question that governance analysis typically evades: how do you get from here to there?

Three structural observations clarify what transition can and cannot accomplish.

Which reforms are architecturally capable

The compounding analysis has a direct implication for reform strategy. Not all reforms are equally capable of reducing the coordination failure tax, even when they are politically feasible and institutionally well-designed.

Reforms that operate within the existing signal architecture — better representatives, more resources, stronger enforcement, cleaner elections, more transparent bureaucracy — cannot address structural constraints. They improve the quality of the governance process acting on the signal after it arrives. If the signal has been destroyed before it arrives — by aggregation, by chain depth, by observation latency — improving what acts on it cannot recover what was lost. Institutional quality is necessary for governance to perform up to its architectural ceiling; it cannot raise the ceiling itself.

Reforms that change the architecture — shortening representation chains, adding observation dimensions, distributing authority to scales appropriate to the disturbance frequencies involved, embedding governance within the systems it governs — operate on the structural constraints directly. They raise the ceiling. The improvements they produce are not incremental improvements to existing performance; they are expansions of the performance envelope.

This distinction is practically significant. It identifies which reform proposals are category errors — improvements to parametric quality within an architecture that cannot perform the required function regardless of parametric quality — and which are genuine structural advances. Much of the energy in contemporary democratic reform focuses on the former: campaign finance, electoral rules, anti-corruption measures, transparency requirements. These are not worthless; they prevent the system from performing below its architectural ceiling. But they cannot address what creates the ceiling.

Partial improvement is real improvement

The coordination failure tax is a continuous function, not a binary one. A governance system need not satisfy all four structural constraints fully to see performance improvement. Because the compounding mechanism means that each failure mode amplifies the others, reducing any one failure mode reduces the amplification it provides to the others.

A system that shortens its representation chain from five layers to two — moving from below the SNR threshold to above it — does not merely improve preference transmission. It reduces the corrupted input that the fractal architecture was receiving, making the frequency coverage more effective. It reduces the phantom preference signal that was authorizing unsustainable commons extraction, making the observational adequacy of commons governance more consequential. One architectural improvement propagates across all four dimensions through the compounding relation.

This has an important implication for transition strategy: reforms do not need to be complete to be worthwhile, and incremental architectural improvement produces compounding returns rather than linear ones. Moving from Architecture B to Architecture C is not halfway to Architecture E. It is a disproportionate fraction of the distance, because the failure modes it addresses stop amplifying the ones that remain.

Phased implementation and proof of concept

The GGF's implementation strategy — beginning with willing bioregions, demonstrating measured performance improvement, and allowing architectural change to achieve legitimacy through demonstrated results rather than political imposition — is consistent with what transition theory predicts for complex institutional change.

Wholesale replacement of existing governance architecture is neither politically feasible nor epistemically warranted. The architectural constraints the series identifies are derived from formal models with simplified assumptions. Real governance systems are more complex. Real communities have specific histories, existing institutions, and legitimate attachments to arrangements that function reasonably well in some dimensions even when they fail structurally in others.

The proof-of-concept approach addresses both problems. It does not require political consensus for wholesale change; it requires only sufficient willingness at the margin to test alternative architectures. And it generates the empirical data — real latency measurements, real preference-policy correlations, real commons collapse rates under different governance architectures — that would transform the framework from formal model to diagnostic tool.

The measurement protocol follows directly from the simulation parameters. Governance latency is measurable as mean time between crisis onset and policy response across a sample of documented cases — for example, the interval between scientific identification of a fisheries stock decline and the implementation of an adjusted quota, or between the first recorded spike in urban housing cost indices and the enactment of planning reform. Preference-policy correlation is measurable by comparing stated preferences with policy outcomes across a sample of decisions. Observation dimensionality is measurable by counting the independent signal dimensions used in commons monitoring. Each parameter that the simulations specify as determining architectural performance is, in principle, measurable from administrative data in real governance systems.

If bioregional pilots produce measurable reductions in governance latency, improvements in preference-policy correlation, and reductions in commons collapse risk — consistent with what the framework predicts — the framework becomes more than a formal model. It becomes a diagnostic instrument with demonstrated predictive validity. That validity is what creates the political conditions for broader architectural change: not argument, but demonstrated performance difference.

V. What remains outside the framework

The engineering framework has a boundary. Acknowledging it is not a concession to critics; it is a condition of intellectual honesty that the framework requires by its own standards.

Legitimacy and consent cannot be derived from engineering reasoning. A governance architecture that is structurally optimal — that satisfies all four constraints, minimizes the coordination failure tax, and demonstrably outperforms alternatives on measured governance outcomes — is not thereby legitimate. Legitimacy requires something the framework cannot supply: the consent of the governed, expressed through processes that the governed themselves regard as authoritative.

This creates a paradox the framework can name but not resolve. The structural findings establish a necessary condition for democratic legitimacy: a governance system that cannot transmit citizen preferences to the policy level, for structural reasons that no institutional improvement can fix, cannot be legitimate in any democratic sense regardless of its formal procedures or popular acceptance. Paper III's finding is not merely a technical observation about signal fidelity — it is a claim about which systems can, in principle, be what they claim to be. A system operating below the SNR threshold cannot reliably transmit the distribution of citizen preferences to the policy layer. The signal that arrives is statistically indistinguishable from noise.

But establishing a necessary condition for legitimacy is not the same as establishing a sufficient one. The framework identifies what must be true for democratic governance to be possible. It does not identify what makes any particular architecture legitimate in the eyes of the people it would govern. That is a political question — resolved through political processes, cultural negotiation, historical reckoning, and the slow work of building trust — that engineering reasoning cannot settle and should not pretend to. The fact that the framework cannot resolve legitimacy questions does not mean those questions are secondary. They are simply outside its scope — and naming that boundary is more honest than either ignoring legitimacy or pretending the framework can address it.

The appropriate stance is not to ignore legitimacy but to be precise about what the framework contributes to it. The framework establishes that systems failing the structural tests cannot be legitimate in a democratic sense. It does not establish that systems passing the structural tests are legitimate. The engineering diagnosis is a floor, not a ceiling. Satisfying the structural constraints is necessary for democratic governance; it is far from sufficient.

There is a further limitation. The framework models governance as a control problem — a system that observes an environment, processes the signal, and produces responses. Real governance also involves narrative, meaning, identity, and the human need to understand oneself as part of a collective project with moral significance. These dimensions are not epiphenomenal to governance; they are often the primary

determinant of whether governance arrangements are accepted, maintained, and extended. The framework has nothing to say about them directly, and should not be read as implying that they can be replaced by architectural optimization.

VI. Conclusion

The Governance as Engineering series began with a simple observation: governance is a form of control, and control has structural prerequisites. A controller that cannot observe the system it governs cannot govern it, regardless of institutional quality. Observation has architecture — channels, latency, dimensionality — and that architecture determines what governance can achieve before institutional quality has anything to act on.

Papers I through IV demonstrated this across four governance domains. Each paper applied a different formal framework — control theory, frequency analysis, information theory, Ashby's Law — to a different problem and arrived at the same structural result: aggregation destroys information; destroyed information cannot be recovered downstream; institutional quality operates on the signal after it arrives and cannot help if the signal is gone.

This paper has made explicit what the series implies but never states: the four failure modes compound. A governance system exhibiting all four simultaneously does not perform four times worse than a structurally sound one. The effective governance capacity is the product of what each failure leaves intact — which, when all four are present at realistic magnitudes, is a small fraction of the baseline. The coordination failure tax is not additive. It is exponential in the number of structural failures present.

The compounding analysis has two practical implications. First, it explains why institutional reforms within the existing Westphalian architecture tend to disappoint: they improve parametric quality within a structure that imposes architectural constraints on performance that no parametric improvement can lift. The ceiling is architectural, not institutional. Second, it means that partial architectural improvement produces disproportionate returns: reducing any one failure mode reduces the amplification it provides to the others, so the gain from architectural improvement exceeds the gain that would be predicted by treating the failure modes as independent.

The Global Governance Frameworks represent one architecture designed to satisfy all four structural constraints simultaneously. Each design principle is load-bearing against a specific failure mode: bioregional positioning against spatial blindness, polycentric nesting against frequency gaps, shallow participation chains against preference invisibility, the Indigenous Wisdom Council and Commons Trusts against observational inadequacy. The architecture is integrated in the specific sense that removing any element leaves the corresponding failure mode unaddressed and restores the coordination failure tax for that dimension.

The GGF is not the only possible architecture satisfying these constraints. The constraints are the point, not the specific design. Any governance architecture that embeds observation within the systems it governs, covers the full disturbance frequency spectrum, transmits citizen preferences through chains short enough to preserve signal, and maintains observation dimensionality matching resource system variety will outperform the Westphalian baseline — measurably, predictably, and for structural rather than contingent reasons.

Measurably and predictably: this is the series' core commitment. The findings are not interpretive claims about which governance systems are philosophically preferable. They are structural claims about what governance systems can achieve under specified architectural conditions. Those claims generate testable predictions — about latency and response quality, about preference-policy correlation and chain depth, about commons collapse rates and observation dimensionality — that empirical work could confirm or falsify.

The framework is an invitation to that empirical work. The simulations establish existence proofs and generate hypotheses. The next step is measurement: real latency distributions in real governance systems, real preference-policy correlations under real chain depths, real resource outcomes under real observation architectures. If the predictions hold, the framework becomes a diagnostic instrument. If they do not, the framework requires revision. Either outcome advances the work.

The coordination failure tax is being paid, continuously and invisibly, by every system that operates below requisite variety across multiple architectural dimensions simultaneously. The series has established what the tax is, why it compounds, and what an architecture that eliminates it would need to look like. Whether that architecture is built is, in the end, a political question. But it is no longer an architectural mystery.