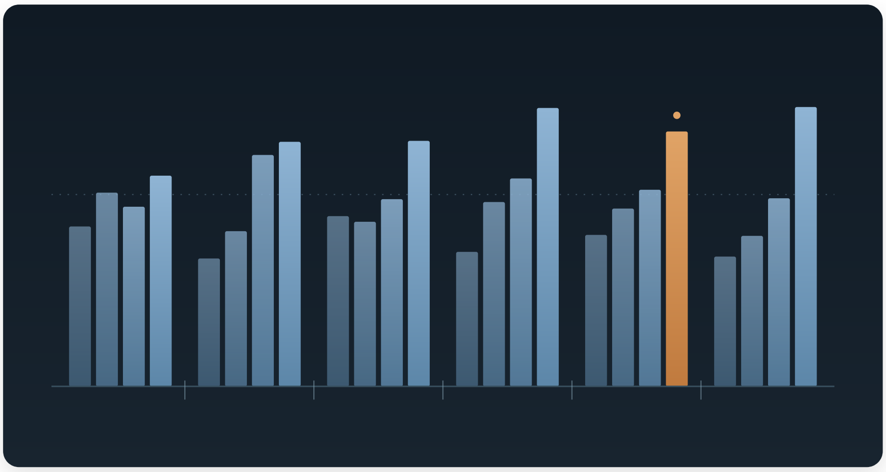

The Four-Year Ramp

Legislative production and the electoral clock in the Swedish Riksdag, 2002–2026

Legislative volume in the Riksdag rises 73% from a mandate's first riksmöte to its fourth, accounting for 82.4% of all variance in output. Which parties govern accounts for 1.4%. Latency is flat at 71–72 days regardless of load, the additional propositions are not smaller, and rejection rates are constant across 23 of 24 riksmöten. The end-of-term surge is well documented; its invariance to government identity is not.

Björn Kenneth Holmström

August 2026

Data: Sveriges riksdag, öppna data (fri användning, ange källa)

`https://bjornkennethholmstrom.org/working-papers/four-year-ramp`

Draft 2 — all figures from Riksdagens öppna data, 24 riksmöten, 5,715 propositions.

Abstract

Swedish election campaigns are argued in part over governing capacity: which constellation can legislate, at what tempo, with what force. This paper tests those claims against the full open-data record of the Riksdag from 2002 to 2026 — six governments, four constitutional arrangements, 5,715 government propositions with chamber outcomes at förslagspunkt resolution.

Legislative volume is governed almost entirely by position in the electoral cycle. Output rises from a mean of 155 non-budget propositions in a mandate's first riksmöte to 268 in its fourth, a 73% increase accounting for 82.4% of all variance. Which parties govern accounts for 1.4%. The year-four surge is not queue-clearing: the additional propositions are no smaller, and the chamber absorbs them with no increase in processing time.

The end-of-term surge itself is well documented; the contribution here is that it does not vary with which parties govern, which is difficult to reconcile with the vote-seeking mechanism the literature ascribes to it.

Two of the four production measures have demonstrated sensitivity — each resolves a large effect elsewhere in the same data — so their failure to detect party effects is informative rather than a power failure. The single detectable deviation in twenty-four years is riksmöte 2021/22, when the government's proposition rejection rate rose fivefold during the collapse of the Löfven III / Andersson government. That result survives four separate tests for coverage artefacts.

The paper does not claim that governments are interchangeable. It claims something narrower and more awkward: on the process measures the Riksdag records, the calendar and acute governmental crisis are visible, and party control is not.

---

1. Introduction

Swedish election campaigns routinely turn on claims about governing capacity. *Handlingskraft*, *reformtempo*, *regeringsduglighet* — the vocabulary presumes that different constellations would produce different amounts of legislation at different speeds with different success. These are empirical claims about the production characteristics of a legislature, and they are testable.

They have not, so far as I can establish, been tested against the full open-data record. The Riksdag publishes every document with structured metadata: submission dates, committee referrals, chamber decisions, and — in the *dokforslag* block — the outcome of every individual förslagspunkt, coded as bifall, delvis bifall, or avslag. Twenty-four riksmöten of this material covers six governments spanning nearly every parliamentary arrangement modern Sweden has produced.

Legislative volume follows a four-year ramp — a pattern already established in the political legislation cycle literature, and expected in a PR system. What is not established, and what this paper reports, is that the ramp is essentially identical under every government observed, that the additional output is not cheaper or faster-processed, and that rejection rates are constant except during one riksmöte of governmental collapse. The variation that campaigns contest does not appear in the record of what parliament produces; the variation that does appear is a property of the calendar.

I want to be careful about what this does and does not establish, and §7 and §9 are devoted to that. The composition question — whether governments differ in *what* they legislate about, as opposed to how much — could not be resolved by this design, and I report the attempt as a limitation rather than a null.

1.1 Related work

Political legislation cycles

The pattern documented in §4 is not new. A literature in political economy, running under the heading *political legislation cycles* (PLC), has established that legislative output rises toward the end of an electoral term. Lagona and Padovano (2008) set out the theory and the first tests; Padovano and Sy (2026) provide the first panel test, covering twenty electoral democracies from the mid-1970s onward, and find that the number of laws falls at the start of a legislature and rises significantly near its end, typically in the final six months. Country-level studies replicate the pattern in the Czech Republic (Brechler and Geršl 2014), in semi-presidential France (Gavoille and Padovano 2017), and for decree production in Italy (Dattilo and Padovano 2025).

Two features of that literature bear directly on the Swedish case. First, Padovano and Sy find cycle magnitude conditional on institutional design: PR systems produce cycles around 67% larger than majoritarian ones. Sweden is PR, and the ratio observed here (1.73 from first to fourth riksmöte) is broadly consistent with that expectation. The present paper therefore does not claim to discover the cycle. It claims that the cycle is invariant to government identity in a way the standard mechanism does not predict.

Second, and more consequentially, PLC theory attributes the pre-election peak to *opportunistic* behaviour: incumbents adopt vote-maximising strategies as elections approach, and the resulting cycle is read as evidence of dynamic inefficiency in the agency relationship between voters and politicians. That interpretation makes predictions this paper can test, because the existing evidence is cross-country while the variation here is within-country across six governments under a constant constitutional framework.

Three of those predictions fail:

1. **Magnitude should vary with electoral incentive.** A government facing likely defeat, or a fragile coalition needing to demonstrate output, has more to gain from a pre-election surge than a secure one. Across six Swedish governments — including one parliamentary majority, one caretaker period, and one that collapsed mid-mandate — mandate identity explains 1.4% of variance in output (§4).
2. **Pre-election legislation should be cheaper.** Vote-seeking output predicts more numerous, smaller, more visible instruments. Year-four propositions are marginally *larger* than year-one propositions on förslagspunkter (§4.1).
3. **A deliberate surge should strain the processing pipeline.** Latency is flat at 71–72 days across all four year-positions, and the 90th percentile falls in year four (§4.2).

An administrative account fits the same data more comfortably: if the utredning–remiss–beredning pipeline has a characteristic duration, work initiated early in a mandate matures late in it, and the ramp is a transit property of the pipeline rather than a strategy of the government moving through it. This paper cannot adjudicate between the two accounts — §10 sets out the SOU-lag test that would — but the invariance result narrows the space the opportunistic account has to occupy.

Government type and legislative performance

A separate literature asks whether minority and coalition governments legislate less effectively than majority ones. Strøm (1984, 1990) established that minority cabinets are rational solutions rather than pathologies, and that they are not straightforwardly less successful. Krauss (2018) provides the large-N comparative test — 197 governments across 21 parliamentary democracies — and finds that minority cabinets do not generally underperform: what matters is whether legislative support is *formalised*, with formal minority cabinets operating in practice much like executive coalitions. Bergman, Ilonszki and others have examined how coalition agreements affect policy productivity, finding effects conditional on cabinet type.

This is the literature the Swedish case speaks to most directly. The period covers a single-party minority with informal support, a majority coalition, a minority coalition, two different formalised cross-bloc arrangements (Decemberöverenskommelsen, Januariavtalet), and a written support contract with staff embedded in Regeringskansliet (Tidöavtalet). Krauss's account predicts that the formalised arrangements should perform comparably to the majority government. They do — but so does everything else, including the informally supported single-party minority. The distinction the literature draws between formal and substantive minority support does not register in Swedish legislative output.

The one exception is instructive and is consistent with Krauss's mechanism read at higher resolution: the rejection-rate spike of 2021/22 (§6) occurs precisely when a formalised arrangement had ceased to function.

Composition over the cycle

On what governments legislate about, as opposed to how much, the relevant work concerns intra-term dynamics rather than partisan difference. Bäck and colleagues (2023) show German MPs' legislative priorities shifting across the electoral cycle, from party-agenda issues early to competitive differentiation late. Work on coalition governments' strategic timing finds the middle of the term used as a window for controversial legislation, where the distance from the next election reduces electoral risk. Both imply that *when* within a term is a live variable for composition — which the design in §7 does not model, and which may be part of why bloc effects were not recoverable there.

The Swedish case in this literature

Sweden appears in the comparative panels but has not, so far as I can establish, been examined at the resolution the open-data record permits: förslagspunkt-level chamber outcomes across every government proposition for twenty-four years. The contribution here is that resolution and the within-country design it supports, not the discovery of the cycle.

2. The Swedish case

The period 2002–2026 supplies unusually varied parliamentary conditions within a constant constitutional framework:

Mandate	Arrangement
Persson II (2002–06)	S single-party minority, V and MP support
Reinfeldt I (2006–10)	Four-party majority coalition
Reinfeldt II (2010–14)	Four-party minority coalition
Löfven I (2014–18)	S+MP minority, Decemberöverenskommelsen
Löfven II/III + Andersson (2018–22)	S+MP minority, Januariavtalet, then collapse
Kristersson (2022–26)	M+KD+L minority, Tidöavtalet with SD

The period includes one genuine parliamentary majority, one four-month caretaker government, one pandemic, three prime ministers inside a single mandate, and two formal cross-party governing contracts of different design. If parliamentary arrangement affects legislative production, the variation required to detect it is present.

3. Data and construction

All data come from Riksdagens öppna data via the `dokumentlista` and `dokumentstatus` endpoints. The unit of time is the riksmöte; mandates group four riksmöten each.

Four measures:

1. **Volume** — government propositions per riksmöte.
2. **Latency** — days from proposition date to chamber decision, taken from the handling betänkande where the decision is not recorded on the proposition itself.
3. **Instrument size** — förslagspunkter per proposition.
4. **Success** — share of propositions with at least one förslagspunkt rejected by the chamber.

3.1 Population statement

Measures 2–4 are computed on **non-budget propositions that reached a betänkande with a förslag block**: 3,658 of 5,715 cached propositions. This restriction is not incidental and should not be read as a sample of all propositions.

Budget propositions are decided by rambeslut rather than punkt by punkt, and consequently carry förslag blocks at 17.1% against 73.1% for ordinary propositions (Fisher OR = 0.08, $p < 1e\text{-}200$). The budget family is also registered as one document per utgiftsområde, so a single budget proposition appears in the record as roughly twenty-seven documents — and the number so registered shifted from about 42 per year to about 11 during 2017/18–2020/21 through a change in titling convention, not in budgeting. Budget documents are therefore excluded from all measures throughout, and volume results are reported both ways.

3.2 A measure discarded

Partial approval (*delvis bifall*) falls monotonically from 18.1% of propositions in 2002–06 to 2.7% in 2022–26. This is not a finding. Mean förslagspunkter per proposition rises over the same period, and the two correlate at Spearman −0.75 across the twenty-four riksmöten: registrars progressively stopped coding partial approval and began splitting punkter instead. The measure records registration practice and is dropped.

4. The electoral ramp

The existence of an end-of-term surge is expected (§1.1). What follows concerns its magnitude, its invariance, and its composition.

Decomposing proposition counts across the 6 × 4 grid of mandates and year-positions:

Source	SS	df	Share
Mandate	735.7	5	**1.4%**
Year-in-mandate	42,450.5	3	**82.4%**
Interaction	8,314.8	15	16.1%

Year-in-mandate means, non-budget propositions: **155** → **182** → **211** → **268**, a ratio of 1.73. Mandate means span 227 to 245 — eight percent, across every arrangement listed in §2.

Non-budget propositions per riksmöte:

Mandate	Y1	Y2	Y3	Y4
Persson II	170	206	191	224
Reinfeldt I	136	165	246	260
Reinfeldt II	181	175	199	261
Löfven I	143	196	221	296
Löfven II/III + Andersson	161	189	209	271
Kristersson	138	160	200	297

4.1 The surge is not queue-clearing

The obvious deflationary reading is that year-four output is thin: minor amendments and technical instruments cleared from a backlog before dissolution. It is not. Median förslagspunkter per proposition by year-position runs 1, 1, 1, 2, and the mean runs 3.0, 2.9, 3.1, 3.7. Year-four propositions are marginally *larger*, so the ramp represents roughly double the legislative content of a mandate's first year on any weighting.

4.2 The chamber absorbs it without delay

Median latency by year-position: **72, 72, 72, 71 days**. The 90th percentile *falls* in year four, from 155 to 140 days.

This is the most mechanically informative result in the paper. A 73% increase in arrival rate with no increase in processing delay is not what a capacity-constrained pipeline does; elementary queueing predicts the opposite. Two readings are available. Either utilisation sits far below one for three years out of four, or beredning latency is not a queue at all but a fixed structural interval — motionstid, remissrunda, utskottsberedning — whose duration is set by procedure rather than by load.

Both readings place the binding constraint outside parliament. The Riksdag is not what limits Swedish legislative output.

4.3 Amplitude: an artefact, not a trend

Cycle amplitude appears to grow across the series when all propositions are included: year-4/year-1 ratios of 1.20, 1.72, 1.32, 1.72, 1.89, 1.92, with a log-scale linear-by-linear interaction at $F(1,14) = 8.28$, $p = 0.012$.

Excluding budget documents, it does not. The ratios become non-monotone (1.32, 1.91, 1.44, 2.07, 1.68, 2.15), Tukey's non-additivity test returns $p = 0.557$, the linear-by-linear term $p = 0.063$, and a permutation test on the six ratios $p = 0.18$.

Budget documents are precisely where the registration conventions of §3.1 shifted. The apparent amplitude trend is therefore most plausibly registrar behaviour rather than legislative behaviour, and it is not claimed here. §10 retains it as a conjecture with a testable implication.

A note on the design: with one observation per cell, a saturated interaction consumes all fifteen residual degrees of freedom. The quantity labelled "interaction" above *is* the interaction; there is no independent error term, and conventional interaction F-tests do not exist in this layout. The one-degree-of-freedom tests above are the available substitutes.

5. Instrument sensitivity

Four null results invite the objection that the instruments are blunt. Rather than validate against an external measure, I pair each null with a large effect the *same* measure detects in the *same* units.

Volume resolves the year-in-mandate effect at $F(3,15) = 25.5$, 82.4% of variance. A measure with that resolution is not blind. Against it, mandate returns $F = 0.27$ — meaning between-mandate variation is *smaller* than the within-mandate residual. This is a statement about effect size, not a failure of power: an F below one indicates that grouping by mandate explains less than an arbitrary partition of the same data would be expected to.

Success rate resolves riksmöte 2021/22 at OR = 4.55, Fisher $p = 3.7\text{e-}05$, while the remaining twenty-three riksmöten are homogeneous to $\chi^2 = 13.4$ on 22 df, **$p = 0.92$**. The measure finds a real deviation where one exists.

Latency and instrument size have no such validation. They are reported descriptively in §4; no null is claimed for them.

6. What happened in 2021/22

Proposition rejection rates by mandate (≥1 punkt rejected, Wilson 95% CI):

Mandate	k/n	Rate	CI
Persson II	7/570	1.23%	0.60–2.51
Reinfeldt I	8/643	1.24%	0.63–2.44
Reinfeldt II	9/575	1.57%	0.83–2.95
Löfven I	10/641	1.56%	0.85–2.85
Löfven II/III + Andersson	22/636	3.46%	2.30–5.18
Kristersson	4/593	0.67%	0.26–1.72

The overall χ^2 across mandates is significant (17.7, df 5, $p = 0.003$), and it is entirely one cell. Löfven II against the other five pooled: OR = 2.81, Fisher $p = 0.0004$. Remove that mandate and the remaining five are indistinguishable: $\chi^2 = 2.54$, df 4, $p = 0.64$. Kristersson against the other four: $p = 0.22$ — nothing.

But the mandate is the wrong unit. Within Löfven II the rejections run 3, 3, 3, 13 across the four riksmöten. The final year alone gives 5.94% against 1.37% for every other riksmöte in the series (OR = 4.55, $p = 3.7e\text{-}05$). Excluding it, Löfven II falls to 2.16% against 1.26%, $p = 0.17$.

2021/22 is the riksmöte in which Löfven resigned, Magdalena Andersson was elected prime minister and resigned the same day following the Miljöpartiet exit, was re-elected, and governed on a budget written by the opposition. A government losing votes at four times the ordinary rate under those conditions is not a finding about governing arrangements. It is a description of a government coming apart.

6.1 The coverage threat, tested four ways

Because förslag-block coverage varies across the series (59.8%–70.7% by mandate) and correlates with rejection rate at the mandate level (rho = 0.71, n = 6), the 2021/22 result could in principle be an artefact of that riksmöte being well covered. Four tests:

1. **Riksmöte-level regression.** Across all 24 riksmöten the coverage–rejection correlation is **+0.04** — essentially zero. The mandate-level rho of 0.71 was an artefact of aggregating twenty-four observations to six. 2021/22 remains the extreme value on the coverage-adjusted residual at $z = +3.31$.
2. **Coverage-matched peers.** Against the thirteen riksmöten within ±3pp of its coverage: 4.41% vs 0.85%, OR = 5.38, $p = 0.0003$.
3. **Mantel–Haenszel by coverage tercile.** OR = 8.71, $p < 0.0001$.
4. **Tipping-point bounds.** For coverage to explain the result, the 1,302 *unobserved* propositions in the other riksmöten would need a rejection rate of **13.41%** — 15.3 times their observed rate of 0.88%. Counting every uncovered proposition as undefeated everywhere, the most pessimistic possible recount, still leaves OR = 5.43, $p = 1.7e\text{-}04$.

The fourth test is the one that settles it. The first three can only fail to find a confound; the fourth states how implausible the confound would have to be.

7. Composition: unresolved

If governments do not differ in how much they legislate, they might still differ in what they legislate about. I could not resolve this, and report the attempt rather than a null.

Because mandates are contiguous blocks of calendar time, any secular drift in policy attention registers as a "mandate effect" with no party content. The test must therefore be on **bloc**, which alternates S–M–S–M and is not collinear with time, after removing a linear trend from each series.

Representation	Bloc variance share	Median p	Time share
Committee (utskott) shares	5.2%	0.167	13.6%
Title topics, TF-IDF	3.7%	0.417	19.4%
Title topics, KB-SBERT	3.0%	0.729	14.7%

Committee shares were restricted to the twelve utskott present across all 24 riksmöten; LU was abolished in 2006 and CU created in its place, so their apparent bloc effects were existence windows correlating with bloc. Title topics had digits stripped, since year tokens otherwise cluster documents by riksmöte, and 708 recurring administrative propositions were excluded by a rule stated in advance — treaty ratifications, annual reports, deployment renewals — leaving 4,178 of 4,894.

Bloc variance shares are stable across k = 10–40 and five clustering seeds, ranging 3.1–4.8% with standard deviations under 1.5pp. The associated p-values are not stable: they range from 0.08 to 1.00 across the same sweep. That instability is itself the clearest illustration of the design's limit. Bloc is constant within a riksmöte, so effective n is 24 regardless of 4,178 documents; the circular-shift null admits 23 labellings, **the minimum achievable p is 0.042**, and the statistic therefore takes a handful of discrete values that move sharply under trivial changes to the partition while the effect size does not.

Why era and party cannot be separated here is visible in the topics themselves. Under both vectorisations the highest bloc-variance cluster is health care and covid — +1.45pp under S-led governments in the SBERT solution — which is the pandemic falling inside an S-led window. Committee data show the same pattern: JuU runs +3.15pp under S-led, coinciding with the gang-violence escalation of 2015–2022, but the M-led government of 2022–26 legislated heavily on crime as well.

Time explains 14–19% of composition variance against bloc's 3–5%. Content appears to track the era more than the government, but that is an observation, not a tested claim. Two unrelated representations agreeing at 3–4% suggests the limitation is the design rather than the vectoriser.

8. Threats to validity

Coverage. Addressed at §6.1 for the 2021/22 result. The riksmöte-level correlation of +0.04 suggests coverage is close to random with respect to contentiousness, but the mechanism generating missing förslag blocks is not established.

Registration practice. Three convention changes were detected and handled: budget document counts (§3.1), partial-approval coding (§3.2), and year tokens in titles (§7). Others may remain undetected. Any longitudinal claim from this source needs the same scrutiny — the amplitude result of §4.3 is a worked example of a finding that did not survive it.

Right-censoring. Riksmöte 2025/26 was incomplete at extraction. Volume may rise further; latency is biased low. The Kristersson rejection rate is stable across three specifications (0.66%, 0.53%, 0.67%), so censoring is not driving it.

Effective sample size. Every mandate- or bloc-level claim runs on 24 observations. Where a result depends on that layer, this is stated.

Selection on submitted propositions. See §9.

9. What this does not show

Nothing about quality or outcomes. A constant rate of legislation says nothing about whether the legislation is good, or whether it achieves anything. Throughput is not effect.

Nothing about composition. §7 is inconclusive in both directions.

Nothing about the current government specifically. The Kristersson mandate is unremarkable on every validated measure. That is a finding about the measures, not a verdict on the government.

Anticipatory adaptation is structurally invisible, and this is the deepest limitation. Governments do not submit propositions they expect to lose. A government skilled at anticipating defeat and a government with a secure majority produce identical rejection rates by opposite routes. What the record contains is the residual after adaptation: bills never drafted, provisions softened in negotiation with a support party, ideas killed at the departementspromemoria stage leave no trace in any document the Riksdag publishes.

This is a general property of institutional measurement rather than a quirk of this dataset. The system records the failures that occurred and is constitutionally blind to those avoided by not attempting them. Any evaluation built on the observable record inherits that blindness, including this one.

10. Implications

For evaluation. Grading a mandate on reformtempo measures the calendar. Any comparison of governments on legislative output that does not condition on year-in-mandate is measuring the electoral clock and attributing it to parties.

For institutional design. If throughput and latency are set upstream of parliament — in Regeringskansliet's utredning capacity, the remiss cycle, and committee procedure — then reform aimed at legislative capacity has to address that apparatus. Nothing in this data suggests parliament is the constraint.

The test that would settle the mechanism. §1.1 set out two competing accounts of the end-of-term surge: the opportunistic one the political legislation cycle literature assumes, in which governments release legislation before elections to maximise votes, and an administrative one in which the ramp is a transit property of a pipeline with a characteristic duration. The invariance results above sit awkwardly with the first but do not establish the second. One measurement discriminates between them. If the ramp is pipeline transit, propositions arriving in year four were initiated earlier and should show systematically longer lags from SOU to proposition than those arriving in year one. If it is strategic release, the lag distribution should be similar across year-positions and the difference should lie in when completed work is published rather than when it was begun. The lag is extractable from the same source — propositions cite the SOU they descend from — and is the natural next step for this line of work.

For the study of political attention. Elections are contested over a variable — which parties govern — that accounts for 1.4% of the variance in how much the legislature produces, and for none of the variance in how fast or how successfully. Whether it accounts for what the legislature produces could not be determined here. If party control matters, this record does not show where.

Appendix: reproduction

All analysis runs from Riksdagens öppna data with no other source. Pipeline: document listing and status extraction, budget-family exclusion, förslagspunkt outcome normalisation, and the tests reported above. Scripts, derived tables, and the API response cache as of 2026-08-18 are available at https://github.com/BjornKennethHolmstrom/riksdag-legislative-production (https://github.com/BjornKennethHolmstrom/riksdag-legislative-production).

Sources: Sveriges riksdag, öppna data (fri användning, ange källa).